Company culture rating — instructions for raters

Round 1 · ground-truth calibration set · issued 2026-08-15

You have six files. Each one is a self-contained web page: the evidence about one company, and the scoring form for that company, in the same window. Open it in a browser. Nothing is installed and nothing is uploaded — your answers stay in the page until you submit.

Budget **about 40 minutes per company**. Six companies, so roughly four hours total. Work at whatever pace suits you — there is no rule about how many you do in a day. The one thing that matters is that you **finish one company before opening the next**, and take a proper break in between. The reading is dense, and the sixth packet of a long day is not being read the same way as the first.

What you are actually doing

You are describing what it is like to work somewhere, on six scales, using only the evidence in front of you.

You are **not** deciding whether the company is good, whether you would work there, or whether the people sound happy. Every one of the six scales runs between two legitimate ways of working. A company at one end is not better than a company at the other end. If you notice yourself scoring toward the end you personally prefer, stop — that is your preference showing up in the data, not the company.

The single most useful thing you can do is keep asking: **which line in this packet made me pick this number?** You will be asked to point at it, every time.

The rules that matter

1. Do not search for anything. Not the company, not a phrase, not a quote. The page asks you at the end whether you searched. Answer honestly — a “yes” costs us almost nothing, and a quiet “no” that isn’t true costs us the whole exercise.

2. If you work out which company it is, that is fine and expected. These packets have had names, links and locations removed, but they have not been made unrecognisable and we are not pretending otherwise. Tell the page you recognised it, say what gave it away, and then keep scoring **from the evidence in front of you rather than from what you know about them.** The recognition answer is data we want, not a confession.

3. One company at a time, and no going back. Each page locks each dimension when you commit it, and locks entirely when you submit. This is deliberate. We want six independent readings, not one profile you have tidied up into a coherent story. Never open two packets side by side, and don’t compare the companies against each other — the numbers are positions on a fixed scale, not a ranking of these six. If you find yourself thinking “this one is more structured than the last one,” that is the thing to avoid; score it against the scale, not against the company you just read.

4. Work alone. Two other people are scoring these same six companies. **Do not discuss the packets, your scores, or your reasoning with them** — not before, not during, not after, not even casually. The entire purpose of using three readers is to find out where three independent judgments diverge. Three people who have

compared notes produce one judgment with three names on it, and there is no way to detect that afterwards or correct for it. Please also don’t forward the files to anyone.

5. Score every dimension, even the ones the packet barely touches. That is what the confidence answer is for. A low-confidence score is useful; a skipped one is not.

6. Read the whole packet before you score anything. The dimensions can be scored in any order once you start.

The two numbers you give for each dimension

The score — five bands: 0.1, 0.3, 0.5, 0.7, 0.9. The page shows you what each band means for that particular dimension. Pick the band whose description best fits the evidence. Do not try to be more precise than the bands allow; we deliberately removed the finer options.

The confidence — High, Medium, Low, Very low. This is about **the evidence, not about you.** “Low” means the packet gave you very little to work with on this dimension. It does not mean you are a hesitant person or that you doubt your own judgment. A confident reader working from a thin packet should still say Low.

Citations — the part that makes this work

Before you can commit a dimension, you must **select at least one passage from the packet** on the page and tag it to that dimension. You select the text; you never type a quote. This guarantees the evidence behind every number is real text from the packet, and it lets us check later that the score actually rests on what it claims to rest on.

You can tag one passage to more than one dimension, and you can attach several passages to a single dimension. Attach the ones that genuinely moved you.

Two things the page does automatically when you select text:

- **It tidies up the edges of your selection** so a citation is always whole words. You do not need to drag precisely.
- **It records which excerpt the passage came from and what kind of source that is** — company-written, an employee review, independent reporting, and so on. It shows you this under the quote before you file it. You do not need to include the source line in your selection; just highlight the words that mattered.

If you highlight something outside the bulleted excerpts, or select only a word or two, the page will tell you and ask for a longer selection from an excerpt.

If you cannot find a passage that supports the number you were about to give, that is a signal. Either the number is coming from somewhere outside the packet, or your confidence should be Very low.

What the packets contain, and what to make of it

Each packet has up to six sections: company careers/values copy, employee reviews, job postings, first-person accounts, recent news, and a human-capital section from the annual report. Some companies are missing some sections entirely — that is a real property of what could be found, not an oversight.

Two things worth knowing as you read:

- **Company-written material tells you what a company says about itself.** Reviews and outside reporting tell you what people experienced. When they disagree, both are evidence, and the disagreement is itself informative.
- **Several packets carry notes about their own limitations** — reviews cut off partway, sections drawn from a narrow slice of the workforce, first-person accounts that are all senior people. Those notes are there because they should change how much weight you put on that section. Read them.

The order you work in

Each rater takes the six companies in a different order, so that whatever drift creeps in over four hours does not land on the same companies for all three of you.

	1st	2nd	3rd	4th	5th	6th
Rater A	P-01	P-02	P-03	P-04	P-05	P-06
Rater B	P-06	P-05	P-04	P-03	P-02	P-01
Rater C	P-04	P-05	P-06	P-01	P-02	P-03

Your files are the ones with your letter in the name — `P-03_raterB.html` and so on. Please work in the order above, not in numerical order.

What happens to your answers

Your six sets of scores are compared against the other two raters', dimension by dimension. Where two of you land in the same band or one band apart, the scale is doing its job. Where you land two or more bands apart, **that is a signal the wording of that anchor is ambiguous, and the wording gets rewritten.** We are not averaging you together and we are not scoring you against each other. Disagreement is the output we are looking for — it is how we find the parts of the scale that are unclear.

The page also records how long you spent, the order you committed the dimensions in, and the version of the anchor wording you saw, so that a later change to the wording can be told apart from a change in judgment.

Getting your answers back to us

When you submit, the page **automatically saves a small file to your Downloads folder** called something like `solverah_P-03_raterB.json`. That file is the only thing we need — **send it to Sebastian.** The web page itself never changes and never needs to be sent back.

Three things worth knowing:

- **Send each file as you finish it**, rather than saving all six for the end.
- **Don’t clear your browser history or site data** until Sebastian has confirmed he has your files. Until you send them, your answers live only on your own machine.
- If the automatic save doesn’t happen for any reason, there are **Download result file** and **Copy to clipboard** buttons on the same screen. Copy-and-paste into an email works fine.

If a red banner appears at the top of the page, your browser is refusing to save progress as you go. You can still score normally and the result file will still work — but finish that packet in one sitting without refreshing, and mention the banner when you send the file.

If something goes wrong

Do not refresh a page mid-company if you can avoid it. If you lose one, tell us which packet and we will reissue it — a re-rated packet is fine as long as we know it happened, and much better than a guessed reconstruction.

If you reopen a packet you already submitted, it stays locked and shows your original result; you can re-download the same file from there if the first one went astray.

If a passage is unreadable, a section seems to contradict itself, or an anchor description does not seem to fit anything in the packet, note it and carry on. Those notes are useful. The instrument is round 1 and it is expected to have rough edges.